

# wt\_petase\_model\_creation

May 6, 2025

## 1 Creating a competitive supervised model for wild type PETase activity at low pH

We show in a separate paper TODO that supervised models should be used for predicting PETase activity at unique conditions (eg. low pH) once some assay labeled data is available, and this outperforms HMMs. Here we create models that: 1. Take in embeddings as input, explore over: Aligned OHE, ES|M2, SaProt, MSATransformer 2. Use linear vs non-linear models: Linear regression, Random Forest

Hyperparameter optimization is conducted over the models for each input type.

Save the final model, which can be loaded like any other sklearn model if AIDE is installed.

eg. `model=joblib.load('model.pkl')`

```
[2]: import os

import pandas as pd
import numpy as np
from sklearn.model_selection import RandomizedSearchCV, KFold, cross_validate
from sklearn.feature_selection import VarianceThreshold
from sklearn.linear_model import ElasticNet, LinearRegression, Ridge
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.pipeline import Pipeline
from sklearn.metrics import roc_auc_score
from sklearn.ensemble import RandomForestRegressor
from scipy.stats import loguniform, spearmanr
import seaborn as sns
import matplotlib.pyplot as plt
sns.set_style("white")
sns.set_context("talk")

import aide_predict as ap
from aide_predict.utils.data_structures.structures import StructureMapper
```

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/patches\_.py:7:

FutureWarning: In the future `np.str` will be defined as the corresponding NumPy

```

scalar.
    if not hasattr(np, 'str'):
/projects/proteinml/.links/miniconda3/envs/aidep/lib/python3.10/site-
packages/Bio/pairwise2.py:278: BiopythonDeprecationWarning: Bio.pairwise2 has
been deprecated, and we intend to remove it in a future release of Biopython. As
an alternative, please consider using Bio.Align.PairwiseAligner as a
replacement, and contact the Biopython developers if you still need the
Bio.pairwise2 module.
    warnings.warn(

```

## 1.1 1. Load and prepare data

We need to get: 1. The sequences and labels 2. Assign their structures (for SaProt Embedding) 3. Get and MSA of known PETases (for baseline HMMscore and MSA transformer)

```
[3]: RAW_DATA_DIR = os.path.join('.', 'data', 'p740')
```

### 1.1.1 1.1 Label data

```
[4]: df = pd.read_csv(os.path.join(RAW_DATA_DIR, 'label_data.csv')).sample(frac=1,
↳ random_state=42)
```

```
[5]: df
```

```
[5]:
```

	Unnamed: 0		sequence	round	\
9	DP021	ADNPYQRGPAPTAASISADTGPFATATTVAEGTGFGGATIYYPTDT...	1		
197	ESM041	AAAAGRADQRGPDPSVAGVAATYGPFATAQLTVPAGNGFNGGYIYY...	3		
66	TEP081	MHPTPDRAKVLPVNVSRGPAEPPAARSARPGGRSAPDGLRPGRRRP...	2		
191	ESM053	VQIGPAPTKASLEASRGPFVATTRLSSANGHGGGTIYYPTNAGAKV...	3		
117	TEP182	MAENPYERGPAPTTSSIEASRGSFATSTVTVSRLAVSGFGGGTIYY...	2		
..	...		...	...	
106	TEP014	ANPYERGPNTQALLEARSGPFSVSSERAWRLGSDGFGGGTIYYPR...	2		
14	DP009	SAQVTRQAAGSYARGPAPTLGIRAALGPFAYSTVTVTAAQAGGAF...	1		
92	TEP024	RPASAQDNPYERGPAPT VSSVAAQRGT FATAELTVPPGNGFNGGKI...	2		
179	ESM011	DSPLYQRGPDPTLASVAATRGPFATTQATVPAGNGFNGGFVYYPTDT...	3		
102	TEP188	ADNPYERGPAPTNASIEAVRGPYAVSQATVSSLAVTGFGGGTIYYP...	2		

	temporal_split	cross_val_split	has_nonzero_activity_anywhere	\
9	0	3	False	
197	4	2	False	
66	1	1	False	
191	4	2	True	
117	2	0	True	
..	...	...	...	
106	2	4	True	
14	0	2	True	
92	2	1	True	
179	4	1	False	

102	2	0	False
-----	---	---	-------

	max_observed_activity	activity_at_5.5_40_cryPow	\
9	0.000	0.00	
197	0.000	0.00	
66	0.000	0.00	
191	0.030	0.01	
117	0.020	0.01	
..	...	...	
106	0.005	0.00	
14	0.010	0.00	
92	0.010	0.00	
179	0.000	0.00	
102	0.000	0.00	

	activity_at_5.5_60_cryPow	activity_at_6.5_40_aFilm	\
9	0.0	NaN	
197	NaN	NaN	
66	NaN	NaN	
191	0.0	NaN	
117	0.0	0.0	
..	...	...	
106	0.0	0.0	
14	0.0	NaN	
92	0.0	0.0	
179	0.0	0.0	
102	0.0	NaN	

	activity_at_6.5_40_cryPow	activity_at_6.5_60_cryPow	\
9	0.00	0.00	
197	0.00	NaN	
66	0.00	NaN	
191	0.01	0.01	
117	0.01	0.01	
..	...	...	
106	0.00	0.00	
14	0.00	0.00	
92	0.01	0.00	
179	0.00	0.00	
102	0.00	0.00	

	activity_at_7.5_40_aFilm	activity_at_7.5_40_cryPow	\
9	NaN	0.000	
197	NaN	0.000	
66	NaN	0.000	
191	NaN	0.030	
117	0.0	0.020	

..	...	...
106	0.0	0.005
14	NaN	0.010
92	0.0	0.010
179	0.0	0.000
102	NaN	0.000

	activity_at_7.5_60_aFilm	activity_at_7.5_60_cryPow \
9	NaN	0.00
197	NaN	NaN
66	NaN	NaN
191	NaN	0.01
117	0.0	0.02
..	...	...
106	0.0	0.00
14	NaN	0.00
92	0.0	0.00
179	0.0	0.00
102	NaN	0.00

	activity_at_8.5_40_cryPow	activity_at_4.5_40_cryPow
9	0.00	NaN
197	NaN	0.0
66	NaN	0.0
191	NaN	0.0
117	NaN	0.0
..	...	...
106	NaN	0.0
14	0.01	NaN
92	NaN	0.0
179	NaN	0.0
102	NaN	0.0


[213 rows x 18 columns]

```
[6]: df = df.dropna(subset=['activity_at_5.5_40_cryPow'])
      # drop rows with non canonical AAs - these were not predicted properly by AF
      X = ap.ProteinSequences.from_df(df, seq_col='sequence', id_col='Unnamed: 0')
      has_non_canonical = [x.has_non_canonical for x in X]
      df = df[~np.array(has_non_canonical)]
```

```
[7]: X, y = ap.ProteinSequences.from_df(df, seq_col='sequence', id_col='Unnamed: 0',
      ↪ label_cols=['activity_at_5.5_40_cryPow'])
```

```
[8]: sns.kdeplot(y, bw_adjust=0.5)
      plt.xlabel('Activity at 5.5 pH, 40°C')
```

```
[8]: Text(0.5, 0, 'Activity at 5.5 pH, 40°C')
```



### 1.1.2 1.2 Structures

```
[9]: # get the structures - Needed for SaProt embedding
mapper = StructureMapper(os.path.join(RAW_DATA_DIR, 'structures'))
mapper.assign_structures(X)
```

```
[9]: ProteinSequences(count=212)
```

### 1.1.3 1.3 Homolog MSA (for MSA transformer)

Compute weights so that MSA transformer can sample it properly.

```
[10]: msa = ap.ProteinSequences.from_fasta(os.path.join(RAW_DATA_DIR, 'hmm-61.mfa'))
msa.aligned
```

```
[10]: True
```

```
[11]: msa.width
```

```
[11]: 898
```

```
[12]: # also assign the msa to the sequences so that msa transformer can access it
      for seq in X:
          seq.msa = msa
```

## 1.2 2. Define scoring functions

```
[13]: # 5 fold cv
      cv_obj = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
[14]: # metrics to measure
      # marks magnitude of error eg R2, also score AUROC to see if the model can
      # ↪ classify active or not
      scoring = {
          'spearman': lambda est, X, y: spearmanr(y, est.predict(X))[0],
          'roc_auc': lambda est, X, y: roc_auc_score(y > 0.001, est.predict(X))
      }
```

```
[15]: def construct_pipeline(embedder, model, pca: bool=True):
      if not pca:
          return Pipeline([
              ('embedder', embedder),
              ('var', VarianceThreshold()),
              ('scaler', StandardScaler()),
              ('model', model)
          ])
      else:
          return Pipeline([
              ('embedder', embedder),
              ('var', VarianceThreshold()),
              ('scaler', StandardScaler()),
              ('pca', PCA(n_components=0.98)),
              ('model', model)
          ])

      def evaluate_pipeline_with_hyperopt(embedder_name, embedder, model_info,
      # ↪ n_iter=10):
          """Run hyperparameter optimization on a pipeline with a given embedder and
          # ↪ model

          Params:
          embedder: Embedder object eg ap.BaseProteinModel
          model_info: dict with keys:
              'model': sklearn model object
```

```

        'param_distributions': dict of hyperparameter distributions for
        ↪RandomizedSearchCV
        """
        do_pca = 'OneHot' not in embedder_name
        pipeline = construct_pipeline(embedder, model_info['model'], pca=do_pca)
        random_search = RandomizedSearchCV(
            pipeline,
            param_distributions=model_info['param_dist'],
            n_iter=n_iter,
            cv=cv_obj,
            scoring=scoring,
            refit='spearman',
            verbose=2,
            n_jobs=1)
        random_search.fit(X, y)

        best_params = random_search.best_params_
        cv_scores = cross_validate(pipeline.set_params(**best_params), X, y,
        ↪cv=cv_obj, scoring=scoring)
        return best_params, cv_scores

```

### 1.3 3. Baseline model: HMM

```

[16]: hmm = ap.HMMWrapper()
      hmm.fit(msa)

      baseline_scores = {
          k: v(hmm, X, y) for k, v in scoring.items()
      }
      print('Baseline scores:', baseline_scores)

```

```

# hmmbuild :: profile HMM construction from multiple sequence alignments
# HMMER 3.4 (Aug 2023); http://hmmer.org/
# Copyright (C) 2023 Howard Hughes Medical Institute.
# Freely distributed under the BSD open source license.
# - - - - -
# input alignment file:
/tmp/HMMWrapper_20250506_092305/alignment.a2m
# output HMM file:
/tmp/HMMWrapper_20250506_092305/alignment.hmm
# - - - - -

# idx name                nseq  alen  mlen  eff_nseq  re/pos  description
#----
1      alignment           61    898   332    4.32    0.590

# CPU time: 0.14u 0.00s 00:00:00.14 Elapsed: 00:00:00.14

```

```
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/predictors/hmm.py:176: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\s+'`` instead
data = pd.read_csv(out_tbl, delim_whitespace=True, comment='#', header=None)
```

```
Baseline scores: {'spearman': -0.06975930257389627, 'roc_auc': 0.47665585919407133}
```

```
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/predictors/hmm.py:176: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\s+'`` instead
data = pd.read_csv(out_tbl, delim_whitespace=True, comment='#', header=None)
```

#### 1.4 4. Supervised learning: Define embedders, models, and hyperparameter space

```
[17]: embedders = {
    'ESM2': ap.ESM2Embedding(
        metadata_folder='esm2_embeddings',
        ↪model_checkpoint='esm2_t33_650M_UR50D', device='cuda:1', pool='mean'),
    'SaProt': ap.SaProtEmbedding(metadata_folder='saprot_embeddings',
        ↪device='cuda:1', pool='mean'),
    'MSATransformer': ap.MSATransformerEmbedding(
        metadata_folder='msa_embeddings', device='cuda:1', pool=False,
        ↪flatten=True, # chosen because there will be a lot of gaps, so mean pool
        ↪will get saturated by gaps
        n_msa_seqs=31, batch_size=32
    ),
    'AlignedOneHot': ap.OneHotAlignedEmbedding(
        metadata_folder='onehot_embeddings')
}
# fit the models that have fixed fitting over folds
embedders['ESM2'].fit()
embedders['SaProt'].fit()
embedders['MSATransformer'].fit()
```

Some weights of EsmModel were not initialized from the model checkpoint at facebook/esm2\_t33\_650M\_UR50D and are newly initialized: ['pooler.dense.bias', 'pooler.dense.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Some weights of EsmModel were not initialized from the model checkpoint at westlake-repl/SaProt\_650M\_AF2 and are newly initialized:

```
['contact_head.regression.bias', 'contact_head.regression.weight',
'embeddings.position_embeddings.weight', 'pooler.dense.bias',
'pooler.dense.weight']
```



You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
[17]: MSATransformerEmbedding(device='cuda:1', flatten=True,
                               metadata_folder='msa_embeddings', n_msa_seqs=31)
```

```
[20]: models = {
    'Ridge': {
        'model': Ridge(),
        'param_dist': {
            'model__alpha': loguniform(1e-5, 1e2),
        }
    },
    'RandomForest': {
        'model': RandomForestRegressor(n_estimators=10),
        'param_dist': {
            'model__max_depth': [None, 10, 100],
            'model__min_samples_split': [2, 5, 10],
            'model__min_samples_leaf': [1, 5, 10]
        }
    }
}
```

## 1.5 5. Train and evaluate models with hyperparameter optimization

```
[21]: import joblib
if not os.path.exists('search_results.pkl'):
    results = {}
else:
    results = joblib.load('search_results.pkl')
```

```
[22]: results
```

```
[22]: {'ESM2_Ridge': {'embedder': 'ESM2',
                    'model': 'Ridge',
                    'best_params': {'model__alpha': 67.83794155101033},
                    'spearman': array([0.73325634, 0.51952198, 0.5185911 , 0.41303091,
0.29284354]),
                    'roc_auc': array([0.8959276 , 0.76973684, 0.76417234, 0.73333333,
0.67108753])},
      'ESM2_RandomForest': {'embedder': 'ESM2',
                           'model': 'RandomForest',
                           'best_params': {'model__min_samples_split': 10,
                                           'model__min_samples_leaf': 5,
                                           'model__max_depth': 10},
                           'spearman': array([0.71394883, 0.52054878, 0.23287475, 0.39442078,
```

```

0.29086657]],
'roc_auc': array([0.89819005, 0.75438596, 0.58956916, 0.72345679,
0.69230769])),
'SaProt_Ridge': {'embedder': 'SaProt',
'model': 'Ridge',
'best_params': {'model__alpha': 86.68878466127286},
'spearman': array([0.71892587, 0.49434321, 0.54815407, 0.44436428,
0.38740863])),
'roc_auc': array([0.89140271, 0.76535088, 0.79365079, 0.74814815,
0.75066313])),
'SaProt_RandomForest': {'embedder': 'SaProt',
'model': 'RandomForest',
'best_params': {'model__min_samples_split': 10,
'model__min_samples_leaf': 10,
'model__max_depth': 10},
'spearman': array([0.68908961, 0.64365332, 0.48380593, 0.50731589,
0.41422383])),
'roc_auc': array([0.8800905 , 0.85526316, 0.75510204, 0.77037037,
0.77320955])),
'MSATransformer_Ridge': {'embedder': 'MSATransformer',
'model': 'Ridge',
'best_params': {'model__alpha': 1.0774036315824223e-05},
'spearman': array([ 0.13592487,  0.21124989,  0.10028396, -0.08393548,
-0.01968863])),
'roc_auc': array([0.5678733 , 0.65131579, 0.60544218, 0.43209877, 0.4668435
])},
'MSATransformer_RandomForest': {'embedder': 'MSATransformer',
'model': 'RandomForest',
'best_params': {'model__min_samples_split': 2,
'model__min_samples_leaf': 10,
'model__max_depth': 100},
'spearman': array([-0.03629287,  0.01247378, -0.1253086 ,  0.00811826,
-0.07255363])),
'roc_auc': array([0.51809955, 0.50877193, 0.4399093 , 0.50493827,
0.40981432])),
'AlignedOneHot_Ridge': {'embedder': 'AlignedOneHot',
'model': 'Ridge',
'best_params': {'model__alpha': 0.001558737330417496},
'spearman': array([0.76998351, 0.71717534, 0.39874602, 0.51453206,
0.47153279])),
'roc_auc': array([0.92986425, 0.90789474, 0.70975057, 0.78765432,
0.78779841])),
'AlignedOneHot_RandomForest': {'embedder': 'AlignedOneHot',
'model': 'RandomForest',
'best_params': {'model__min_samples_split': 2,
'model__min_samples_leaf': 5,
'model__max_depth': 10},

```

```

'spearman': array([0.68039666, 0.52833455, 0.3305744 , 0.40241285,
0.43515628]),
'roc_auc': array([0.86877828, 0.80701754, 0.64172336, 0.71604938,
0.76127321])}]

```

```

[23]: for embedder_name, embedder in embedders.items():
        for model_name, model_info in models.items():
            if f'{embedder_name}_{model_name}' in results:
                print(f"Skipping {embedder_name} with {model_name}...")
                continue
            else:
                print(f"Evaluating {embedder_name} with {model_name}...")
                best_params, scores = evaluate_pipeline_with_hyperopt(embedder_name,
↪embedder, model_info, n_iter=50)
                results[f'{embedder_name}_{model_name}'] = {
                    'embedder': embedder_name,
                    'model': model_name,
                    'best_params': best_params,
                    'spearman': scores['test_spearman'],
                    'roc_auc': scores['test_roc_auc']
                }
                joblib.dump(results, 'search_results.pkl')

```

```

Skipping ESM2 with Ridge...
Skipping ESM2 with RandomForest...
Skipping SaProt with Ridge...
Skipping SaProt with RandomForest...
Skipping MSATransformer with Ridge...
Skipping MSATransformer with RandomForest...
Skipping AlignedOneHot with Ridge...
Skipping AlignedOneHot with RandomForest...

```

```

[24]: # convert to long
df_list = []
for item in results.values():
    for i in range(5): # Assuming 5 values for each metric
        df_list.append({
            'embedder': item['embedder'],
            'model': item['model'],
            'spearman': item['spearman'][i],
            'roc_auc': item['roc_auc'][i]
        })

df = pd.DataFrame(df_list)

# Melt the DataFrame to create a column for the metric type

```

```
df_melted = pd.melt(df, id_vars=['embedder', 'model'], var_name='metric',  
↳ value_name='value')
```


```
[25]: df_melted.to_csv('fig3_data.csv', index=False)
```

```
[30]: plt.figure(figsize=(8, 10))  
  
# Create the faceted plot  
g = sns.catplot(  
    data=df_melted,  
    kind="bar",  
    x="model",  
    y="value",  
    hue="metric",  
    col="embedder",  
    height=4,  
    aspect=1.0,  
    palette="Set2",  
    col_wrap=4,  
    ci="sd",  
    legend=True, # We'll add the legend manually  
    # change color  
    edgecolor='black',  
    linewidth=2  
  
)  
# remove legend from seaborn  
g._legend.remove()  
# add baselines of the same color as the bars  
for ax in g.axes.flat:  
    for i, metric in enumerate(['spearman', 'roc_auc']):  
        ax.axhline(baseline_scores[metric], color=sns.color_palette("Set2")[i],  
↳ linestyle='--', label=f'HMM (baseline) {metric}')  
        ax.set_ylim(-.1, 1)  
  
plt.legend(loc='upper center', bbox_to_anchor=(-1.1, -0.18), ncol=4)  
# Customize the plot  
g.set_axis_labels("")  
g.set_titles("{col_name}")  
  
# lower the  
  
# Display the plot  
plt.savefig('p740_model_comparison.png', bbox_inches='tight', dpi=300)
```

/tmp/ipykernel\_2947768/1272998501.py:4: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar='sd'` for the same effect.

```
g = sns.catplot(
<Figure size 800x1000 with 0 Axes>
```



## 1.6 6. Train final model and save

```
[31]: best_row = df.groupby(['embedder', 'model']).mean().sort_values('roc_auc',
    ↪ascending=False).iloc[0]
best_row
```

```
[31]: spearman    0.574394
      roc_auc     0.824592
      Name: (AlignedOneHot, Ridge), dtype: float64
```

```
[36]: best_pipeline = construct_pipeline(embedders[best_row.name[0]], models[best_row.
    ↪name[1]]['model'], pca=False)
```

```
[37]: best_params = results[f'{best_row.name[0]}_{best_row.name[1]}']['best_params']
```

```
[38]: models[best_row.name[1]]['model']
```

```
[38]: Ridge(alpha=0.001558737330417496)
```

```
[39]: best_pipeline.set_params(**best_params)
```

```
[39]: Pipeline(steps=[('embedder',
    OneHotAlignedEmbedding(metadata_folder='onehot_embeddings')),
    ('var', VarianceThreshold()), ('scaler', StandardScaler()),
    ('model', Ridge(alpha=0.001558737330417496))])
```

## 1.7 First do a CV prediction so we can plot parity

```
[40]: y_trues = []
      y_preds = []
      for train_idx, test_idx in cv_obj.split(X):
          best_pipeline.fit(X[train_idx], y[train_idx])
          y_pred = best_pipeline.predict(X[test_idx])
          y_trues.append(y[test_idx])
          y_preds.append(y_pred)

      y_trues = np.concatenate(y_trues)
      y_preds = np.concatenate(y_preds)
      y_trues = y_trues > 0.001
```

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original

```
alignment.")
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/embedder
s/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the
original alignment.
    warnings.warn("Input sequences are not aligned. Aligning them to the original
alignment.")
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/embedder
s/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the
original alignment.
    warnings.warn("Input sequences are not aligned. Aligning them to the original
alignment.")
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/embedder
s/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the
original alignment.
    warnings.warn("Input sequences are not aligned. Aligning them to the original
alignment.")
```

```
[41]: def Find_Optimal_Cutoff(target, predicted):
    """ Find the optimal probability cutoff point for a classification model
    ↳related to event rate
    Parameters
    -----
    target : Matrix with dependent or target data, where rows are observations

    predicted : Matrix with predicted data, where rows are observations

    Returns
    -----
    list type, with optimal cutoff value

    """
    from sklearn.metrics import roc_curve
    fpr, tpr, threshold = roc_curve(target, predicted)
    i = np.arange(len(tpr))
    roc = pd.DataFrame({'tf' : pd.Series(tpr-(1-fpr), index=i), 'threshold' :
    ↳pd.Series(threshold, index=i)})
    roc_t = roc.iloc[(roc.tf-0).abs().argsort()[:1]]

    return list(roc_t['threshold'])

cutoff = Find_Optimal_Cutoff(y_trues, y_preds)
cutoff
```

```
[41]: [0.005645442926271432]
```

```
[42]: precision = np.sum((y_preds > cutoff) & y_trues) / np.sum(y_preds > cutoff)
precision
```

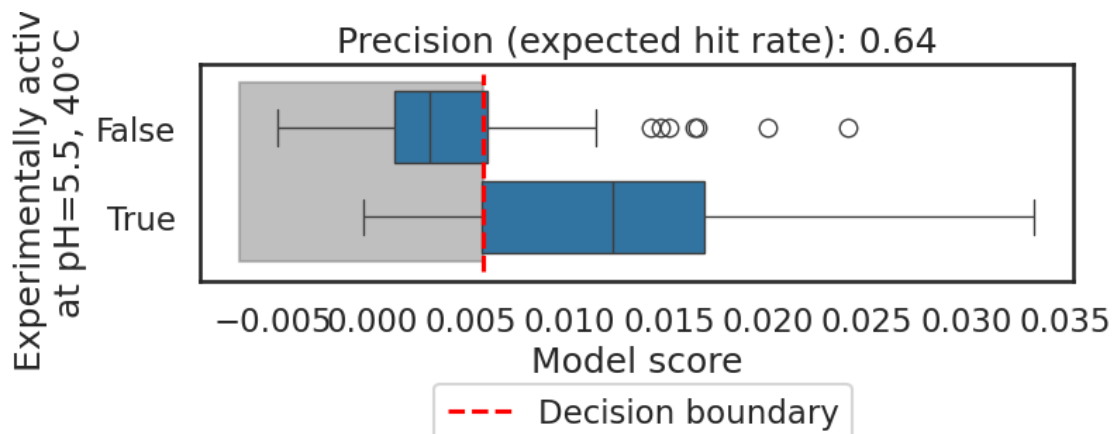
[42]: 0.6391752577319587

```
[45]: fig, ax = plt.subplots(figsize=(8, 2))

df_ = pd.DataFrame({
    'y_true': y_trues.flatten(),
    'y_pred': y_preds.flatten(),
})

sns.boxplot(data=df_, x='y_pred', y='y_true', ax=ax, orient='h')
ax.set_xlabel('Model score')
ax.set_ylabel('Experimentally active \nat pH=5.5, 40°C')
ax.set_title('Precision (expected hit rate): {:.2f}'.format(precision))

ax.fill_between([ax.get_xlim()[0], cutoff[0]], [ax.get_ylim()[0], ax.
    ↪get_ylim()[0]], [ax.get_ylim()[1], ax.get_ylim()[1]], color='grey', alpha=0.
    ↪5)
ax.vlines(cutoff[0], ax.get_ylim()[0], ax.get_ylim()[1], color='red',
    ↪linestyle='--', label='Decision boundary')
plt.legend(bbox_to_anchor=(.25, -.4), loc='upper left')
plt.savefig('p740_precision.png', bbox_inches='tight', dpi=300)
```



```
[46]: best_pipeline = best_pipeline.fit(X, y)
```

/kfs2/projects/proteinml/repos/aide\_predict/aide\_predict/bespoke\_models/embedder  
s/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the  
original alignment.

warnings.warn("Input sequences are not aligned. Aligning them to the original  
alignment.")

```
[47]: preds = best_pipeline.predict(X)
```



```
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.
```

```
warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")
```

```
[48]: joblib.dump(best_pipeline, 'p740_best_pipeline.pkl')
```

```
[48]: ['p740_best_pipeline.pkl']
```

## 1.8 7. Load model and predict

```
[49]: import joblib
best_pipeline = joblib.load('p740_best_pipeline.pkl')
```

```
[50]: preds = best_pipeline.predict(X)
```

```
/kfs2/projects/proteinml/repos/aide_predict/aide_predict/bespoke_models/embedders/ohe.py:278: UserWarning: Input sequences are not aligned. Aligning them to the original alignment.
```

```
warnings.warn("Input sequences are not aligned. Aligning them to the original alignment.")
```

```
[51]: preds
```

```
[51]: array([[ 4.09470608e-10],
          [ 2.62595394e-09],
          [-1.49346806e-10],
          [ 9.99999950e-03],
          [ 1.00000015e-02],
          [ 3.05852435e-10],
          [ 4.82003894e-09],
          [ 3.98192960e-10],
          [ 8.21041975e-09],
          [ 6.91821198e-10],
          [ 9.99999856e-03],
          [ 3.99999322e-02],
          [ 2.14445496e-09],
          [ 2.64940886e-10],
          [-5.17693135e-10],
          [ 2.99999893e-02],
          [ 2.99999703e-02],
          [ 1.63759887e-09],
          [ 3.36233957e-09],
          [ 1.99999930e-02],
          [ 1.02378027e-08],
          [-4.85636236e-10],
```

[ 2.28727708e-09],  
[ 1.00000000e-02],  
[ 1.99999995e-02],  
[ 1.99999967e-02],  
[ 1.22329702e-09],  
[ 9.9998230e-03],  
[ 1.81703685e-09],  
[ 4.79542311e-10],  
[ 2.99999674e-02],  
[ 4.9999815e-02],  
[ 2.93682240e-09],  
[ 9.43059512e-10],  
[ 1.81236652e-10],  
[ 3.40490398e-09],  
[-1.66555745e-09],  
[ 9.9999614e-03],  
[ 9.9999811e-03],  
[ 9.9999452e-03],  
[ 1.9999937e-02],  
[ 3.9999899e-02],  
[-8.41008469e-11],  
[ 1.9999952e-02],  
[ 7.51803440e-09],  
[ 1.9999993e-02],  
[ 4.9999910e-03],  
[ 1.9999954e-02],  
[ 4.23487816e-10],  
[ 6.53770181e-10],  
[ 2.90781677e-09],  
[ 5.96691266e-10],  
[ 1.08329407e-09],  
[ 1.9999946e-02],  
[ 1.90299555e-09],  
[ 1.9999986e-02],  
[ 4.84301615e-10],  
[ 2.0000030e-02],  
[ 4.82629546e-09],  
[ 4.4892308e-09],  
[ 9.9999994e-03],  
[ 9.9999343e-03],  
[ 9.9999742e-03],  
[ 9.9999987e-03],  
[ 1.74454894e-08],  
[ 1.0000049e-02],  
[ 8.23468618e-10],  
[ 3.74034980e-10],  
[ 1.81816497e-09],

[ 2.04303712e-09],  
[ 9.99999785e-03],  
[ 9.9998983e-03],  
[ 1.56742151e-09],  
[ 9.9999570e-03],  
[ 3.10711320e-10],  
[ 7.10378426e-09],  
[ 2.78351574e-09],  
[ 5.14522925e-09],  
[ 9.83719645e-10],  
[ 3.11738056e-10],  
[ 6.64873797e-10],  
[ 1.37286022e-09],  
[ 3.78983846e-09],  
[ 9.9999585e-03],  
[ 5.24034515e-09],  
[ 1.71144497e-09],  
[ 9.9999841e-03],  
[ 1.9999834e-02],  
[ 5.88160239e-09],  
[ 1.84952039e-10],  
[ 9.9999969e-03],  
[ 2.9999945e-02],  
[ 9.9999831e-03],  
[ 9.26675082e-11],  
[ 5.87372872e-09],  
[ 1.12435145e-08],  
[ 5.36214040e-10],  
[ 1.87812789e-09],  
[ 1.67306797e-09],  
[ 1.04455376e-09],  
[ 1.38335092e-09],  
[ 9.9998835e-03],  
[-2.70749365e-10],  
[ 2.90283882e-09],  
[ 3.24197599e-09],  
[ 9.9999437e-03],  
[ 1.88300813e-09],  
[ 9.9999646e-03],  
[ 1.9999961e-02],  
[ 1.49155519e-09],  
[ 2.36352568e-09],  
[ 1.16145253e-09],  
[ 3.55474928e-09],  
[ 9.98354955e-10],  
[ 1.66175792e-09],  
[ 1.9999995e-02],

[ 2.61023906e-09],  
[ 1.00000058e-02],  
[ 1.56580710e-09],  
[ 7.68034560e-10],  
[ 1.00000008e-02],  
[ 3.88377377e-09],  
[ 1.50000142e-02],  
[ 1.33091808e-10],  
[ 5.76884456e-09],  
[ 1.00000026e-02],  
[ 9.69297915e-11],  
[ 2.99999717e-02],  
[ 3.89117379e-09],  
[ 2.89131351e-09],  
[ 1.00000263e-02],  
[ 3.37972027e-09],  
[ 3.99999714e-02],  
[ 2.50280752e-09],  
[ 1.04850780e-08],  
[ 7.51749635e-11],  
[ 3.21758458e-09],  
[ 5.19999843e-02],  
[ 2.00000164e-02],  
[ 3.19438413e-09],  
[ 1.49999974e-02],  
[ 9.99999675e-03],  
[ 9.99999703e-03],  
[ 1.99999971e-02],  
[ 9.99999835e-03],  
[ 5.00000568e-03],  
[ 2.48338943e-09],  
[ 1.99999915e-02],  
[ 4.99999863e-02],  
[ 2.49999910e-02],  
[ 1.58545946e-09],  
[ 1.49999967e-02],  
[ 2.99999932e-02],  
[ 5.68301739e-09],  
[ 8.80656181e-09],  
[ 8.50935275e-10],  
[ 5.39803827e-10],  
[ 1.08686036e-09],  
[ 1.00000045e-02],  
[ 9.99999787e-03],  
[ 9.35810171e-09],  
[ 1.00000159e-02],  
[ 7.39228063e-10],

[-2.38702096e-10],  
[ 2.00000117e-02],  
[ 2.60109920e-10],  
[ 1.52954845e-09],  
[ 1.99999961e-02],  
[ 1.68812737e-09],  
[ 1.37579154e-09],  
[ 1.00000162e-02],  
[ 1.50000188e-02],  
[ 1.30482774e-09],  
[ 4.08411322e-09],  
[ 8.65600377e-09],  
[ 9.99999592e-03],  
[ 1.00000001e-02],  
[ 1.49999947e-02],  
[ 1.96691585e-09],  
[ 1.50000117e-02],  
[ 2.44886844e-10],  
[ 1.30798086e-09],  
[ 2.49999974e-02],  
[ 1.07223158e-09],  
[ 1.00000053e-02],  
[ 1.00000095e-02],  
[ 2.16768613e-09],  
[ 1.81267683e-09],  
[ 1.00000188e-02],  
[ 2.01293365e-09],  
[ 1.04728018e-09],  
[ 9.99999461e-03],  
[ 4.88358373e-10],  
[ 5.54614634e-09],  
[ 2.99999822e-02],  
[ 1.76539674e-10],  
[ 1.26165288e-09],  
[ 1.00000070e-02],  
[ 2.96559986e-09],  
[ 1.00000070e-02],  
[ 3.78598194e-09],  
[ 1.04319380e-09],  
[ 9.99999594e-03],  
[ 5.68815614e-10],  
[ 9.66619623e-11],  
[ 2.99999743e-02],  
[ 5.58217745e-10],  
[ 2.08647065e-09],  
[ 3.40143808e-09],  
[ 3.34551228e-09],

```
[ 1.95501513e-09],  
[ 6.85415975e-10]])
```

```
[ ]:
```